

Dynamic KL prior adds small cost at low data,
matches baseline at full training size

