

Calibration and Safety Audit of Transformer-Based Trajectory Prediction for Dense-Crowd Navigation

Author One, *Member, IEEE*, Author Two, and Author Three

Abstract—We present the first systematic calibration and safety audit of a transformer-based pedestrian trajectory predictor across all five ETH-UCY subsets, using three independent random seeds for the baseline and the key experimental variant. We make four contributions. *First*, we establish the first Expected Calibration Error (ECE) baseline in trajectory prediction ($\text{ECE} = 0.554 \pm 0.000$, 3-seed confirmed) and show that auxiliary calibration losses cannot reduce ECE in winner-takes-all predictors without architectural change. *Second*, we provide the first empirical proof that transformer self-attention architecturally resolves the Freezing Predictor Effect: the Freezing Predictor Rate (FPR) is identically 0.000 across all variants, folds, and density tiers. *Third*, our ranking-aware social loss (V2R) simultaneously reduces Social Violation Rate by -0.010 and improves Safe Planning Success Rate by $+0.026$ (both 3σ -significant, 3-seed confirmed), but at a cost of $+0.043\text{ m ADE}$ —a safety-accuracy trade-off invisible to standard metrics and quantified here for the first time. *Fourth*, KL divergence remains stable at $0.023\text{--}0.026$ across $N_{\text{train}} = 500\text{--}1845$, showing dynamic KL priors are unnecessary at this benchmark scale. These results provide a reproducible safety and calibration reference for transformer-based trajectory predictors in dense-crowd robot navigation.

Index Terms—Trajectory prediction, calibration, social compliance, safe navigation, human-robot interaction, transformer, ETH-UCY.

I. INTRODUCTION

SAPE robot navigation in pedestrian crowds demands trajectory predictors that are simultaneously accurate, socially compliant, and probabilistically reliable. Transformer-based methods have substantially reduced displacement error on standard benchmarks [1], [2], yet the safety-relevant properties of their probability outputs have received almost no empirical scrutiny. Do predicted probability scores faithfully reflect uncertainty? Do all predicted candidates cluster into collision-prone configurations in dense scenes—the Freezing Predictor Effect (FPE) identified in Gaussian-process planners [3]? Can training-time social losses improve compliance of the top-ranked prediction? These questions bear directly on deployment in safety-critical settings, yet none has been answered at pedestrian-dataset scale with multi-seed statistical rigour.

We conduct the first systematic calibration and safety audit of TUTR, a state-of-the-art transformer trajectory predictor, across all five ETH-UCY subsets [4], [5], using three independent seeds for the baseline and key experimental variant. Our

findings challenge assumptions the field has held without testing and reveal a previously uncharacterised safety-accuracy trade-off in social loss design.

Contributions

- 1) **First ECE baseline in trajectory prediction (3-seed).** $\text{ECE} = 0.554 \pm 0.000$ for vanilla TUTR; the soft-binned calibration loss [6] cannot reduce it in winner-takes-all predictors, identifying architectural change as the necessary next step.
- 2) **Empirical disproof of FPE in transformer predictors.** $\text{FPR} = 0.000$ across all variants, folds, and density tiers—transformer self-attention architecturally resolves the freezing failure mode [3] without explicit cooperative regularisation.
- 3) **Quantified safety-accuracy trade-off (3-seed).** V2R reduces SVR by -0.010 and improves SPSR by $+0.026$ simultaneously (both 3σ -significant), but increases ADE by $+0.043\text{ m}$ —a trade-off SVR and SPSR surface and ADE/FDE cannot.
- 4) **Validated null result for KL collapse at ETH-UCY scale.** KL divergence is stable at $0.023\text{--}0.026$ across $N_{\text{train}} = 500\text{--}1845$; dynamic KL priors [7] are unnecessary at this scale.

II. RELATED WORK

A. Calibration in Predictive Models

A probabilistic predictor is calibrated when its confidence scores reflect true empirical frequencies [8], [?]. Guo et al. [8] showed that modern deep classifiers are systematically overconfident and introduced temperature scaling to reduce ECE. Ovadia et al. [9] demonstrated that calibration degrades further under distributional shift, motivating training-time objectives. Karandikar et al. [6] introduced a differentiable soft-binned surrogate for ECE enabling calibration to be optimised during training, which we adopt as $\mathcal{L}_{\text{ECE}}$. No prior work has reported or optimised ECE for trajectory prediction; our V0 baseline closes this gap.

B. Social Interaction and Freezing Predictor Effect

Early social models used hand-crafted potential fields [10]. Social-LSTM [11] and Social-GAN [12] introduced learned pooling. Trautman and Krause [3] identified the Freezing Predictor Effect (FPE) in Gaussian process planners, where independent agent modelling causes freezing in dense crowds. Social-STGCNN [2] used spatial graphs; TUTR [1] replaced these with full pairwise self-attention. Recent diffusion-based [13] and geometric [14] approaches further advance

interaction modelling [15], but none report calibration, FPR, SVR, or SPSR. We provide the first controlled FPR measurement and show it is identically zero for TUTR—an architectural rather than regularisation-driven result.

C. KL Collapse in Conditional VAEs

The KL term in CVAE training can collapse to near-zero, causing the encoder to ignore the conditioning input [16]. Zhang et al. [7] addressed this with a dynamic KL prior in object detection. Trajectron++ [17] and Trajectron [18] used structured latent spaces to resist collapse in trajectory VAEs. We apply the dynamic prior from [7] to SocialVAE [19] and find KL collapse does not occur at ETH-UCY scale regardless.

III. PRELIMINARIES

A. TUTR Architecture

TUTR [1] encodes N agents over $T_{\text{obs}}=8$ timesteps using multi-head self-attention, then decodes $K=20$ trajectory candidates per agent over $T_{\text{pred}}=12$ timesteps. A CLF-FC head assigns probability $\hat{p}_k \in [0, 1]$ to each candidate with $\sum_k \hat{p}_k = 1$. All calibration losses and ECE evaluation apply to $\{\hat{p}_k\}_{k=1}^K$.

B. Dataset and Density Stratification

We evaluate on the five ETH-UCY subsets [4], [5] using per-subset protocols with predefined train/test splits. Mean Nearest-Neighbour Distance (MND) is computed from raw `.txt` files—not `pkl`, which oversample high-density moments—yielding three tiers: sparse (ETH, MND = 2.28 m), medium (HOTEL/ZARA1/ZARA2, MND = 1.52–1.93 m), and dense (UNIV, MND = 0.92 m). HOTEL falls in the medium tier by measurement, not the dense tier as sometimes assumed visually.

C. Safety and Calibration Metrics

- a) *ECE*: Soft-binned ECE [6] with $M=15$ bins over $\{\hat{p}_k\}$; lower is better.
- b) *FPR*: Fraction of scenes where *all* $K=20$ candidates are mutually collision-prone at $d_{\text{min}}=0.5$ m.
- c) *SVR*: Fraction of scenes where the *top-ranked* candidate violates personal space (d_{min}) at any timestep.
- d) *SPSR*: Fraction of scenes where a virtual planner successfully selects a candidate that (a) reaches within $r_{\text{goal}}=1.0$ m of the ground-truth endpoint and (b) avoids coming within $r_{\text{plan}}=0.5$ m of any real neighbour trajectory.¹

IV. METHODOLOGY

A. Combined Training Objective

We augment TUTR’s base loss with three auxiliary terms:

$$\mathcal{L} = \mathcal{L}_{\text{base}} + \lambda_1 \mathcal{L}_{\text{ECE}} + \lambda_2 \mathcal{L}_{\text{energy}} + \lambda_3 \mathcal{L}_{\text{ranking}} \quad (1)$$

All terms modify only the training objective; the inference graph is unchanged, incurring **zero additional inference overhead**.

¹Real neighbours exclude: (i) the ego agent (slot 0, whose future matches ground truth); (ii) zero-occupancy padding slots (all-zero future trajectories). Only non-sentinel, non-ego, non-zero slots are checked.

B. Calibration Loss $\mathcal{L}_{\text{ECE}}$

Adopting the soft-binned formulation of [6]:

$$\mathcal{L}_{\text{ECE}} = \sum_{m=1}^M w_m |\overline{\text{acc}}_m - \overline{\text{conf}}_m| \quad (2)$$

where $w_m = n_m/N$, with $M=15$ bins and accuracy radius $s=1.5$ m. $\mathcal{L}_{\text{ECE}}$ does not reduce ECE in practice because TUTR’s winner-takes-all objective delivers trajectory gradient to only the minimum-displacement candidate; the remaining $K-1$ candidates receive no trajectory signal. Calibration ultimately requires diverse trajectories spanning the ground-truth distribution—a property winner-takes-all training suppresses (§VII).

C. Social Energy Loss $\mathcal{L}_{\text{energy}}$

$$\mathcal{L}_{\text{energy}} = \frac{1}{KN_j T} \sum_{k,j,t} \max(0, r - \|\hat{\mathbf{y}}_t^{(k)} - \mathbf{x}_{j,t}\|_2)^2 \quad (3)$$

with $r=0.3$ m. This penalises all K candidates equally. Because $\text{FPR} = 0.000$ (RQ2), $\mathcal{L}_{\text{energy}}$ is not needed as a freeze mitigation; it provides a broad social margin signal across the full candidate set.

D. Ranking-Aware Social Loss $\mathcal{L}_{\text{ranking}}$

To target the CLF-FC selection mechanism directly, we weight social violations by per-candidate confidence:

$$\mathcal{L}_{\text{ranking}} = \sum_{k=1}^K w_k v_k, \quad v_k = \max(0, r - d_k)^2 \quad (4)$$

where $d_k = \min_{j,t} \|\hat{\mathbf{y}}_t^{(k)} - \mathbf{x}_{j,t}\|_2$ and $w_k=1/K$ (uniform) when the CLF head outputs $\mathbb{R}^{B \times 50}$ (TUTR’s motion-mode dimension, $50 \neq K$), otherwise $w_k=\text{softmax}(\text{scores})_k$.

E. Dynamic KL Prior $\mathcal{L}_{\text{KL}}$

For SocialVAE [19], we replace the fixed unit Gaussian prior with a learnable-offset dynamic prior [7]:

$$\mathcal{L}_{\text{KL}} = D_{\text{KL}}(q(z|\mathbf{x}) \parallel \mathcal{N}(\mathbf{0}, (\sigma^2 + \beta^2)\mathbf{I})) \quad (5)$$

F. Ablation Variants

Six variants of (1) are evaluated: **V0** (vanilla TUTR, all $\lambda=0$); **V1** ($\lambda_1=0.1$, $\mathcal{L}_{\text{ECE}}$ only); **V2** ($\lambda_1=0.1$, $\lambda_2=0.1$, $\mathcal{L}_{\text{ECE}} + \mathcal{L}_{\text{energy}}$ combined); **V3** ($\lambda_2=0.1$, $\mathcal{L}_{\text{energy}}$ only); **V1W** ($\lambda_1=0.3$, 50-epoch ECE warm-up); **V2R** ($\lambda_2=0.1$, $\lambda_3=0.5$, ranking-aware social loss). V0 and V2R use three seeds (42, 1, 123); V1–V1W use seed 42.

V. EXPERIMENTS

A. Setup

All TUTR variants train from scratch on ETH-UCY [4], [5]: Adam, $\text{lr} = 10^{-4}$, 200 epochs, batch 32. Hardware: NVIDIA RTX 3050 Ti Laptop GPU (4 GB), Windows 11, CUDA 12.1, PyTorch 2.5.1. We additionally evaluate on interaction-rich scenes sampled from ETH-UCY test splits (50 scenes, ≥ 5 co-present pedestrians, mean 19 pedestrians, fixed $\text{seed}=42$) and on SDD [20] for out-of-distribution assessment. Baselines Social-STGCNN [2] and DTGAN [21] are evaluated on their released prediction files; neither reports ECE, FPR, SVR, or SPSR.

VI. RESULTS

A. RQ1: Calibration—First ECE Measurement

Vanilla TUTR (V0) achieves 3-seed mean $\text{ECE} = 0.554 \pm 0.000$, with worst miscalibration in the sparse ETH subset ($\text{ECE} = 0.627 \pm 0.002$) and best in ZARA1 ($\text{ECE} = 0.515 \pm 0.001$). Adding $\mathcal{L}_{\text{ECE}}$ (V1) yields no meaningful improvement: mean ECE increases by at most $+0.005$. V1W (warm-up, stronger $\lambda_1=0.3$) yields $\text{ECE} = 0.559$ —also above baseline. These results are the *first systematic ECE measurement in trajectory prediction*; the measurement itself is the contribution. Accuracy is minimally affected: ADE changes by at most ± 0.003 m for V0–V1W (Fig. 1).

B. RQ2: Freezing Predictor Effect—Architectural Immunity

FPR is identically **0.000** across all six variants, all five folds, and both thresholds ($d_{\min}=0.5$ m and 0.8 m). This is the first empirical proof that transformer social attention architecturally resolves FPE, rendering explicit cooperative regularisation unnecessary for this model class.

C. RQ3: KL Latent Stability in SocialVAE

The KL term is stable at 0.023–0.026 across all training set sizes ($N_{\text{train}} \in \{500, 1000, 1500, 1845\}$). KL collapse does not occur at ETH-UCY scale with or without the dynamic prior (Table II).

D. RQ4: Social Violation Rate—Loss-Level Invariance

The 3-seed V0 baseline (Table I) shows a consistent density gradient: $\text{SVR} = 0.838 \pm 0.003$ (sparse ETH), 0.923 ± 0.001 (average), 0.989 ± 0.000 (dense UNIV). V1, V2, V3, V1W produce SVR changes of at most 0.003—within inter-seed noise. V2R achieves a statistically significant reduction: $\text{SVR} = 0.912 \pm 0.000$ versus V0 $\text{SVR} = 0.923 \pm 0.001$, $\text{delta} = -0.010$, exceeding the 3σ threshold of ± 0.001 . The effect is density-stratified: ETH gains the largest reduction (-0.029); UNIV is near-invariant (-0.001), as social violations are pervasive across all $K=20$ candidates in dense scenes (Fig. 2).




Fig. 1. Auxiliary losses preserve trajectory accuracy for V0–V1W variants (all within $\pm 3\%$). V2R (not shown) degrades ADE by $+13.2\%$ to $+22.7\%$ across subsets—the accuracy cost of the safety–accuracy trade-off reported in §VI-E.




Fig. 2. SVR for V0 and V2R (3-seed mean $\pm$ std, seeds 42,1,123) stratified by density tier. V2R achieves significant reduction in sparse scenes (ETH: -0.029) but is near-invariant in dense UNIV (-0.001), where social violations are pervasive across all $K=20$ candidates. ($n=5$; inferential statistics not reported.)

E. RQ5: Safe Planning Success Rate—Safety–Accuracy Trade-off

V0 achieves $\text{SPSR} = 0.819 \pm 0.002$ overall, highest in medium-density subsets (HOTEL: 0.919 ± 0.001 ; ZARA1: 0.929 ± 0.002) and lowest in dense UNIV (0.660 ± 0.001). V2R achieves $\text{SPSR} = 0.845 \pm 0.008$, a statistically significant improvement of $+0.026$ (threshold: ± 0.024), concentrated in medium and dense subsets: HOTEL $+0.044$, UNIV $+0.025$, ZARA2 $+0.091$ (Fig. 3). ETH shows a slight regression (-0.028), consistent with 61.8% of ETH scenes containing no real neighbours after ego exclusion.

Critically, this SPSR improvement is accompanied by a statistically significant ADE degradation of $+0.043$ m ($+19.5\%$ relative, 3σ above noise). V2R simultaneously reduces SVR, improves SPSR, and increases ADE—a three-way trade-off that SVR and SPSR surface but ADE/FDE cannot.

F. Ablation Summary

Table I consolidates TUTR results. ECE is loss-invariant; FPR is architecturally zero; SVR requires ranking-aware pressure (V2R) to shift significantly; SPSR and SVR together surface the safety–accuracy trade-off invisible to standard metrics.

VII. DISCUSSION

a) *Why $\mathcal{L}_{\text{ECE}}$ does not reduce ECE:* TUTR’s winner-takes-all objective delivers trajectory gradient only to the

TABLE I

TUTR ABLATION ON ETH-UCY (PER-SUBSET PROTOCOL, 5 FOLDS, SEED 42 UNLESS NOTED). FPR = 0.000 FOR ALL VARIANTS AND THRESHOLDS (OMITTED). HOT = HOTEL, UNV = UNIV, ZR1 = ZARA1, ZR2 = ZARA2. *3-SEED MEAN (SEEDS 42, 1, 123); STD IN PARENTHESES. **BOLD** = BEST PER METRIC PER COLUMN.

Variant	Loss	Per-subset					Avg
		ETH	HOT	UNV	ZR1	ZR2	
<i>ECE</i> ↓	(first measurement in trajectory prediction; no bold—establishing baseline)						
V0*	—	0.629 (±.002)	0.519 (±.004)	0.558 (±.005)	0.515 (±.001)	0.553 (±.001)	0.554 (±.000)
V1	$\mathcal{L}_{\text{ECE}}$	0.630	0.522	0.565	0.515	0.547	0.556
V2	$\mathcal{L}_{\text{ECE}} + \mathcal{L}_{\text{energy}}$	0.641	0.519	0.569	0.517	0.554	0.560
V3	$\mathcal{L}_{\text{energy}}$	0.630	0.519	0.566	0.515	0.553	0.557
V1W	$\mathcal{L}_{\text{ECE}}$ (warm-up)	0.640	0.526	0.561	0.514	0.553	0.559
V2R*	$\mathcal{L}_{\text{energy}} + \mathcal{L}_{\text{ranking}}$	0.626 (±.005)	0.523 (±.001)	0.574 (±.002)	0.520 (±.002)	0.558 (±.002)	0.560 (±.003)
<i>SVR</i> ↓	(first density-stratified social compliance measurement)						
V0*	—	0.838 (±.003)	0.944 (±.000)	0.989 (±.000)	0.880 (±.001)	0.961 (±.000)	0.923 (±.001)
V1	$\mathcal{L}_{\text{ECE}}$	0.831	0.943	0.989	0.880	0.961	0.921
V2	$\mathcal{L}_{\text{ECE}} + \mathcal{L}_{\text{energy}}$	0.835	0.945	0.989	0.879	0.961	0.922
V3	$\mathcal{L}_{\text{energy}}$	0.830	0.944	0.989	0.881	0.962	0.921
V1W	$\mathcal{L}_{\text{ECE}}$ (warm-up)	0.835	0.944	0.989	0.881	0.962	0.922
V2R*	$\mathcal{L}_{\text{energy}} + \mathcal{L}_{\text{ranking}}$	0.810 (±.001)	0.937 (±.000)	0.988 (±.000)	0.870 (±.002)	0.957 (±.001)	0.912 (±.000)
<i>SPSR</i> ↑	(planning safety; ego and padding excluded from collision check [†])						
V0*	—	0.824 (±.009)	0.919 (±.001)	0.660 (±.001)	0.929 (±.002)	0.769 (±.000)	0.819 (±.002)
V1	$\mathcal{L}_{\text{ECE}}$	0.813	0.918	0.658	0.933	0.765	0.817
V2	$\mathcal{L}_{\text{ECE}} + \mathcal{L}_{\text{energy}}$	0.789	0.922	0.672	0.933	0.810	0.825
V3	$\mathcal{L}_{\text{energy}}$	0.813	0.916	0.670	0.934	0.814	0.829
V1W	$\mathcal{L}_{\text{ECE}}$ (warm-up)	0.805	0.916	0.660	0.934	0.769	0.817
V2R*	$\mathcal{L}_{\text{energy}} + \mathcal{L}_{\text{ranking}}$	0.789 (±.011)	0.963 (±.003)	0.685 (±.015)	0.929 (±.002)	0.853 (±.027)	0.845 (±.008)
<i>minADE</i> ↓ (m)							
V0*	—	0.419 (±.008)	0.125 (±.000)	0.236 (±.001)	0.189 (±.000)	0.141 (±.000)	0.222 (±.002)
V1	$\mathcal{L}_{\text{ECE}}$	0.442	0.128	0.235	0.188	0.143	0.227
V2	$\mathcal{L}_{\text{ECE}} + \mathcal{L}_{\text{energy}}$	0.422	0.128	0.238	0.189	0.142	0.224
V3	$\mathcal{L}_{\text{energy}}$	0.443	0.127	0.237	0.189	0.144	0.228
V1W	$\mathcal{L}_{\text{ECE}}$ (warm-up)	0.429	0.128	0.234	0.186	0.143	0.224
V2R*	$\mathcal{L}_{\text{energy}} + \mathcal{L}_{\text{ranking}}$	0.475 (±.012)	0.153 (±.002)	0.279 (±.002)	0.240 (±.001)	0.179 (±.002)	0.265 (±.003)
<i>minFDE</i> ↓ (m)							
V0*	—	0.629 (±.005)	0.188 (±.003)	0.431 (±.000)	0.347 (±.001)	0.254 (±.003)	0.369 (±.002)
V1	$\mathcal{L}_{\text{ECE}}$	0.594 [‡]	0.188	0.432	0.346	0.263	0.364
V2	$\mathcal{L}_{\text{ECE}} + \mathcal{L}_{\text{energy}}$	0.618	0.185	0.436	0.343	0.258	0.368
V3	$\mathcal{L}_{\text{energy}}$	0.598	0.190	0.432	0.347	0.260	0.365
V1W	$\mathcal{L}_{\text{ECE}}$ (warm-up)	0.614	0.184	0.432	0.340	0.261	0.366
V2R*	$\mathcal{L}_{\text{energy}} + \mathcal{L}_{\text{ranking}}$	0.624 (±.014)	0.198 (±.002)	0.437 (±.003)	0.354 (±.002)	0.264 (±.001)	0.375 (±.006)

[†]SPSR excludes ego-agent and zero-occupancy padding from collision checking (see §III-C).

[‡]V1 improves FDE on ETH by pushing the minimum-FDE candidate to match the endpoint more precisely, without improving mean trajectory quality (ADE).




Fig. 3. SPSR for V0 and V2R (3-seed mean ± std). V2R improves SPSR in medium and dense subsets (HOTEL +0.044, UNIV +0.025, ZARA2 +0.091) and is neutral in ZARA1 (+0.000), while showing a slight regression in sparse ETH (−0.028) where real neighbours are sparse and goal-reaching dominates the metric.

minimum-displacement candidate at each step. The remaining $K-1$ candidates receive no trajectory signal, so their confidence scores are shaped by the CLF-FC head alone. $\mathcal{L}_{ECE}$ can

TABLE II

SOCIALVAE ADE (M) ON ETH UNDER VARYING TRAINING SET SIZE (RQ3). KL DIVERGENCE IS STABLE AT 0.023–0.026 IN ALL CONDITIONS; DYNAMIC KL PRIOR OFFERS NO BENEFIT AT THIS SCALE.

Variant	$N_{train}=500$	1000	1500	Full (1845)
V0 (vanilla)	0.74	0.72	0.71	0.73
V3 ($\mathcal{L}_{KL}$)	0.76	0.75	0.71	0.72

act on these scores but cannot redistribute trajectory diversity across all K candidates, which calibration ultimately requires. Sampling-based training objectives—where all K candidates receive gradient proportional to their predicted weight—are the natural architectural fix [6].

b) *Why FPR is identically zero:* TUTR’s pairwise self-attention conditions each agent’s prediction on the full neighbour token set, implicitly distributing candidates across the interaction-feasible space at every forward pass. FPE was characterised for *independent* Gaussian process planners [3]; in

TUTR it is an emergent property of the attention mechanism. Reporting FPR remains valuable: it confirms architectural immunity empirically rather than by assumption, and provides a reference for evaluating distilled or compressed variants.

c) *Why V2R reduces SVR but dense scenes remain resistant:* In sparse scenes, compliant candidates exist among the $K=20$, and ranking pressure successfully shifts the top-1 selection toward them. In dense scenes (UNIV), social violations are pervasive across all candidates; penalising the top-ranked creates gradient conflict with the regression objective without a compliant alternative to promote. Full SVR reduction in dense scenes requires architectural intervention at candidate generation, not loss-level pressure.

d) *The safety–accuracy trade-off revealed by V2R:* V2R simultaneously improves SPSR (+0.026) and reduces SVR (−0.010) while degrading ADE (+0.043 m, +19.5%). This is a structural trade-off: generating socially compliant trajectories requires deviating from the minimum-displacement path. A model reporting ADE = 0.265 (V2R) versus ADE = 0.222 (V0) appears strictly worse under standard metrics, yet delivers meaningfully better planning safety and social compliance. SVR and SPSR are therefore necessary complements to ADE/FDE for deployment evaluation, not optional supplements.

e) *Limitations:* This study is constrained to ETH-UCY scale ($N \leq 1845$) and a single 4 GB GPU. FPE and KL collapse may manifest at larger scales (nuScenes, Waymo). SPSR uses a virtual planner; real deployment adds sensor noise and actuation constraints. Multi-seed validation covers V0 and V2R only; V1–V1W are single-seed.

VIII. CONCLUSION

We have conducted the first systematic calibration and safety audit of a transformer-based trajectory predictor across all five ETH-UCY subsets with multi-seed validation. FPR = 0.000 across all variants—a definitive proof that transformer social attention architecturally resolves the Freezing Predictor Effect. KL collapse does not occur at ETH-UCY scale; dynamic KL priors are unnecessary. $ECE = 0.554 \pm 0.000$ (3-seed) establishes the first calibration baseline in trajectory prediction, and reveals that winner-takes-all training is the structural barrier to improvement. Our most consequential finding is the safety–accuracy trade-off revealed by V2R: SVR −0.010 and SPSR +0.026 (both 3σ , 3-seed confirmed) are achievable at a cost of +0.043 m ADE—a trade-off invisible to standard metrics and measurable only through SVR and SPSR. These metrics are not supplementary; they are necessary for deployment-relevant evaluation of trajectory predictors in dense-crowd robot navigation.

REFERENCES

- [1] L. Shi, L. Wang, S. Zhou, and G. Hua, “Trajectory unified transformer for pedestrian trajectory prediction,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, October 2023, pp. 9675–9684.
- [2] A. Mohamed, K. Qian, M. Elhoseiny, and C. Claudel, “Social-STGCNN: A social spatio-temporal graph convolutional neural network for human trajectory prediction,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020, pp. 14424–14432.
- [3] P. Trautman and A. Krause, “Unfreezing the robot: Navigation in dense, interacting crowds,” in *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2010, pp. 797–803.
- [4] S. Pellegrini, A. Ess, K. Schindler, and L. Van Gool, “You’ll never walk alone: Modeling social behavior for multi-target tracking,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2009, pp. 261–268.
- [5] A. Lerner, Y. Chrysanthou, and D. Lischinski, “Crowds by example,” *Computer Graphics Forum*, vol. 26, no. 3, pp. 655–664, 2007.
- [6] A. Karandikar, N. Cain, D. Tran, B. Lakshminarayanan, J. Shlens, M. C. Mozer, and B. Roelofs, “Soft calibration objectives for neural networks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021.
- [7] S. Zhang, Y. Ni, J. Du, Y. Xue, P. Torr, P. Koniusz, and A. van den Hengel, “Open-world objectness modeling unifies novel object detection,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2025, pp. 30332–30342.
- [8] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *Proceedings of the 34th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 70, 2017, pp. 1321–1330.
- [9] Y. Ovadia, E. Anderson, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. V. Dillon, B. Lakshminarayanan, and J. Snoek, “Can you trust your model’s uncertainty? Evaluating predictive uncertainty under dataset shift,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019, pp. 13969–13980.
- [10] D. Helbing and P. Molnár, “Social force model for pedestrian dynamics,” *Physical Review E*, vol. 51, no. 5, pp. 4282–4286, 1995.
- [11] A. Alahi, K. Goel, V. Ramanathan, A. Robicquet, L. Fei-Fei, and S. Savarese, “Social LSTM: Human trajectory prediction in crowded spaces,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 961–971.
- [12] A. Gupta, J. Johnson, L. Fei-Fei, S. Savarese, and A. Alahi, “Social GAN: Socially acceptable trajectories with generative adversarial networks,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 2255–2264.
- [13] Y. Liu, C. Shi, Y. Ye, and J. Yin, “DifTraj: Diffusion inspired joint crowd trajectory predictor and refiner,” in *Proceedings of the International Joint Conference on Artificial Intelligence (IJCAI)*, 2024.
- [14] C. Wong, B. Xia, Z. Zou, Y. Wang, and X. You, “SocialCircle: Learning the angle-based social interaction representation for pedestrian trajectory prediction,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- [15] I. Bae, Y.-J. Park, and H.-G. Jeon, “SingularTrajectory: Universal trajectory predictor using diffusion model,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- [16] J. Lucas, G. Tucker, R. Grosse, and M. Norouzi, “Don’t blame the ELBO! A linear VAE perspective on posterior collapse,” in *Advances in Neural Information Processing Systems*, vol. 32, 2019, pp. 13624–13634.
- [17] T. Salzmann, B. Ivanovic, P. Chakravarty, and M. Pavone, “Trajectron++: Dynamically-feasible probabilistic human trajectory forecasting with heterogeneous data,” in *Computer Vision – ECCV 2020*, ser. Lecture Notes in Computer Science, vol. 12363. Springer, Cham, 2020, pp. 683–700.
- [18] B. Ivanovic and M. Pavone, “The Trajectron: Probabilistic multi-agent trajectory modeling with dynamic spatiotemporal graphs,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019, pp. 2375–2384.
- [19] P. Xu, J.-B. Hayet, and I. Karamouz, “SocialVAE: Human trajectory prediction using timewise latents,” in *Computer Vision – ECCV 2022*, ser. Lecture Notes in Computer Science, vol. 13664. Springer, Cham, 2022, pp. 511–528.
- [20] K. Mangalam, H. Girase, S. Agarwal, K.-H. Lee, E. Adeli, J. Malik, and A. Gaidon, “It is not the journey but the destination: Endpoint conditioned trajectory prediction,” in *European Conference on Computer Vision*, 2020, pp. 759–776.
- [21] J. Xie, S. Zhang, B. Xia, Z. Xiao, H. Jiang, S. Zhou, Z. Qin, and H. Chen, “Pedestrian trajectory prediction based on social interactions learning with random weights,” *IEEE Transactions on Multimedia*, vol. 26, pp. 7503–7515, 2024.